

Paresh Prajapati

AI/ML Engineer | Generative AI Specialist

+91 8849581032 | prajapatipareshkumar1032@gmail.com | [LinkedIn](#) | [GitHub](#)

SUMMARY

AI/ML Engineer with 3+ years of experience taking Generative AI products from design to production. Shipped three live AI platforms — a catering marketplace assistant, a compliance intelligence engine, and a travel companion — built on agentic architectures (LangChain/LangGraph), RAG pipelines, and AWS AI services. Focused on solutions that balance model quality with cost efficiency, latency, and data privacy.

TECHNICAL SKILLS

GenAI & Agents: LLMs, RAG, LangChain, LangGraph, Agentic AI, Prompt Engineering, NLP, Vector Search

Languages & Backends: Python, FastAPI, Django, Flask, REST APIs, SSE Streaming

Cloud & Data: AWS (Bedrock, ECS Fargate, Textract), OpenSearch, PostgreSQL, Redis, Docker

ML Libraries: Pandas, NumPy, TensorFlow, Scikit-learn, Hugging Face Transformers

Tools & Practices: Git/GitHub, OOP, DBMS, strict typing (mypy), linting & pre-commit quality gates

EXPERIENCE

AI Engineer

Sept 2025 – Present

- Own end-to-end AI solution development — system design, agent architecture, deployment, and monitoring on AWS.
- Design scalable multi-agent workflows (LangGraph) powering conversational assistants, semantic search, and document intelligence in production.
- Drive cost efficiency across LLM workloads through model selection, caching strategies, token-usage tracking, and streaming architectures.

Software Engineer (AI/ML)

Feb 2023 – Oct 2025

- Built and shipped AI/ML solutions optimizing operational processes — RAG frameworks, LLM integrations, and NLP pipelines.
- Collaborated across all project stages, from problem definition and data preparation to model evaluation and delivery.
- Developed production chatbots, document-analysis systems, and recommendation engines used by real clients.

FEATURED PRODUCTION PROJECTS

NG Cater — AI-Powered Catering Marketplace

In Production

- **Business:** built the AI layer of a catering marketplace — customers assemble orders conversationally in plain English, and caterers digitize PDF/photo menus into structured, searchable listings in minutes instead of hours of manual entry.
- **Agent stack:** LangGraph multi-step agent on AWS Bedrock (Claude) with tool-calling into hybrid keyword + vector semantic search over OpenSearch.
- **Production hardening:** persistent conversation memory via LangGraph PostgreSQL checkpointing, PII-safe SSE token streaming, Redis caching, containerized services on ECS Fargate.
- **Menu extraction:** AWS Textract OCR + LLM structuring pipeline converting raw menu documents into normalized catalog data.

Ctrl Shft — Compliance & Policy Intelligence Engine

In Production

- **Business:** automated compliance workflows for regulated care-sector organizations — policy drafting/rewriting aligned to regulatory frameworks, incident and complaint triage, excursion risk assessment, provider document verification, and CV parsing — cutting hours of manual review to minutes while improving auditability.
- **Architecture:** suite of independent AI microservices (FastAPI), each with a strict JSON contract for safe backend integration.
- **Grounding & governance:** RAG over ingested regulatory frameworks so outputs cite the rules they follow; schema-validated structured outputs; per-request AI usage and cost tracking.

My Greek Odyssey — AI Travel Companion

In Production

- **Business:** AI concierge for a Greece travel platform — answers destination, itinerary, and logistics questions from curated knowledge, turning static content into an interactive planning experience.
- **Backend:** FastAPI service with clean route → controller → service architecture, agentic chat with knowledge-base retrieval, and background workers.
- **Production hygiene:** per-request tracing, structured logging, uniform response envelopes, centralized error handling, strict typing (mypy --strict), Docker; consumed by a Flutter cross-platform client.

ADDITIONAL PROJECTS

- **Database Communication Chatbot:** LLM converts natural-language questions into SQL executed via PrestoDB over Cassandra, opening complex datasets to non-technical users.
- **Legal IntelliSearch:** RAG framework for legal document analysis — precise extraction with citations, in-document highlighting, and clause comparison across contracts.
- **Intelligent HR Assistant:** voice-enabled agentic assistant with dynamic tool triggering for leave balances, holiday lookups, and RAG-grounded policy answers.
- **Affinity AI:** LLM-based matchmaking recommendations with demographic rules; early-stopping cut unnecessary API spend and parallel batch execution processed thousands of profiles efficiently.

EDUCATION

B.Tech, Information & Communication Technology (ICT) — Sankalchand Patel	2019 – 2023
University, Visnagar CGPA: 8.33	
Class XII — Gurkul School of Science, Siddhpur	2018 – 2019
Class X — SMT S.B Doshi High School, Nandotra	2016 – 2017

LANGUAGES

English (Proficient) • Hindi (Proficient) • Gujarati (Native)